



Planning with Temporal Logic

April 25, 2016



Motivation

- Consider a self-driving car...
- Regardless of our destination, we also want to make sure we always follow the rules of the road.






Motivation





Motivation





Key Takeaways

- Modeling temporally-extended goals with linear temporal logic (LTL)
- Modeling preferences between alternative plans



Outline

- **Introduction to Linear Temporal Logic**
 - Why use Linear Temporal Logic?
 - Linear Temporal Logic Operators
 - Example LTL Problems
- **Applications to Planning**
- **Planning with Preferences**
 - Expressing Preferences
 - Planning in LPP



Linear Temporal Logic





Temporal Logic

- Formalism for specifying properties of systems that vary with time



Temporal Logic


- Systems proceed through a sequence of discrete states






Why Temporal Logic?

- Previously our planning algorithms have used propositional logic to specify goals dealing with a **single state** at a single point in time




- Temporal logic allows these goals to be specified over a sequence of states





Why Temporal Logic?


- What if the problem requires a condition to:
 - Be met until another condition is met...
 - For example: **red** implies (**stop** until **green**)





Why Temporal Logic?

- What if the problem requires a condition to:
 - Always eventually be met
 - For example, always have some point in the future when you visit a gas station





Branching vs linear time

- Linear time
 - Models physical time
 - At each time instant, only one of the future behaviors is considered
 - We can reason about **always**






Temporal Logic

Branching vs linear time

- Branching time
 - At each time instant, all possible future behaviors are considered
 - Time may split into alternate courses
 - We can reason about **possibilities**






Temporal Logic

Branching vs linear time

- Linear time



- Branching time






Linear Temporal Logic



- Linear Temporal Logic (LTL) involves:

- Linear time model
- Infinite sequences of states



- Forward-looking conditions

- Cannot express properties over a set of different paths





Applications of Temporal Logic

- Temporal logic is used in:
 - Verification and Model Checking
 - Safety and Maintenance
 - **Planning**



LTL Syntax

LTL formula $f := \text{true} \mid p_i \mid f_i \wedge f_j \mid \neg f_i \mid X f_i \mid f_i U f_j$

An LTL formula is built from:

1. **Propositional variables:** p, ρ, ϕ, ω etc.

— Can be True or False

2. **Logical Operators:** $\neg, \vee, \wedge, \rightarrow, \leftrightarrow, \text{True}, \text{False}$

— \neg = not

— \vee = or

— \wedge = and

— \rightarrow = implies

— \leftrightarrow = if and only if

— True, False




Logical Operator Examples

Logical Operators

true

Example

true




Logical Operators

$p = \text{true}$

Example

R = red light






Logical Operator Examples

Logical Operators

not, \neg

Example

$\neg G = \text{green light}$




Logical Operators

and, \wedge

Example

$R \wedge B = \text{gas station}$






Logical Operator Examples

Logical Operators

Example

or, \vee

R \vee **G**



Or (\vee) can be rewritten with and (\wedge) and not (\neg)

$$\mathbf{R} \vee \mathbf{G} = \neg(\neg \mathbf{R} \wedge \neg \mathbf{G})$$

Similar process can be done for implies and iff, but we won't be explaining them due to time constraints



LTL Syntax

LTL formula $f := \text{true} \mid p_i \mid f_i \wedge f_j \mid \neg f_i \mid X f_i \mid f_i U f_j$

An LTL formula is built from:

1. **Propositional variables:** p, ρ, ϕ, ω etc.

— Can be True or False

2. **Logical Operators:** $\neg, \vee, \wedge, \rightarrow, \leftrightarrow, \text{True}, \text{False}$

— \neg = not

— \vee = or

— \wedge = and

— \rightarrow = implies

— \leftrightarrow = if and only if


— True, False

3. Temporal Operators



Temporal Operators


What are some useful operators we may want to describe our car?






Temporal Operators


- The **next** light to be **green**



- The light will be **red** **until** it is **green**




- The light will **eventually** some point in the **future**, turn **green**






Temporal Operators

- The light will **always** be **red**



- The light will be **red** until the car gets **gas** and the state after it's **released**, the light can be whatever





Next


Operator

Textual Operator

ne**X**t

X ρ

Definition: Variable ρ must be true in the next state





Until


Operator

Textual Operator

Until

$\rho \mathbf{U} \omega$

Definition: Variable ρ must remain true up until the state where variable ω becomes true, at which point ρ becomes unconstrained



Note that ω is required to become true in some future state



Future


Operator

Textual Operator

Future/Eventually

Fp

Definition: Variable p must become true in some future state





Global


Operator

Textual Operator

Globally

$G\rho$

Definition: Variable p must be true in all future states





Release


Operator

Textual Operator

Release

$\rho R \omega$

Definition: Variable ρ must be true up until and including the state where ω becomes true, after which ω is unconstrained. If ρ is not true in any future state, then ω is true in all future states



Different from **U** in that both ρ and ω are true in one state




Which describe the other?

Future/Eventually

Release

Globally



?

?

?

$$\equiv \text{True} \cup p$$

$$\equiv \neg F \neg p$$

$$\equiv \neg(\neg p \cup \neg \omega)$$



Globally






Temporal Operators (Recap)

| Operator | Textual Operator |
|---------------------------|--|
| ne X t | X ρ |
| U ntil | ρ U ω |
| F uture/Eventually | F $\rho \equiv \text{True U } \rho$ |
| G lobally | G $\rho \equiv \neg \text{F} \neg \rho$ |
| R elease | ρ R $\omega \equiv \neg(\neg \rho \text{ U } \neg \omega)$ |




Combination of Operators

Infinitely Often




Eventually Forever





Example Problem

What are some true statements about this LTL formation?



- **XR**
- **FG**
- **RUG**
- **(RUG) ∧ (FG) ∧ (XR)**



PDDL3 Goal Description

```
<GD> ::= (at end <GD>)  
      | (always <GD>)  
      | (sometime <GD>)  
      | (within <num> <GD>)  
      | (at-most-once <GD>)  
      | (sometime-after <GD> <GD>)  
      | (sometime-before <GD> <GD>)  
      | (always-within <num> <GD> <GD>)  
      | (hold-during <num> <num> <GD> | ...
```




Temporal Operators

| Operator | PDDL3 |
|-----------------|---|
| neXt | X ρ (within 1 ρ) |
| Until | ρ U ω (always-until ρ ω) |
| Future | ρ F ω (sometime-after ρ ω) |
| Globally | G ρ (always ρ) |
| Release | ρ R ω (or (always ω) (always-until ω ρ)) |



Expressing Temporal Logic in PDDL

- The traffic light will turn red in the **next** state



```
(:goal (within 1 (turn red) ) )
```

Command Syntax

```
(within <num> <GD>)
```


`(within <num> φ)` would mean that φ must hold within
<num> happenings



Expressing Temporal Logic in PDDL

- The traffic light will be green **until** it turns red at which point it will be red **forever**

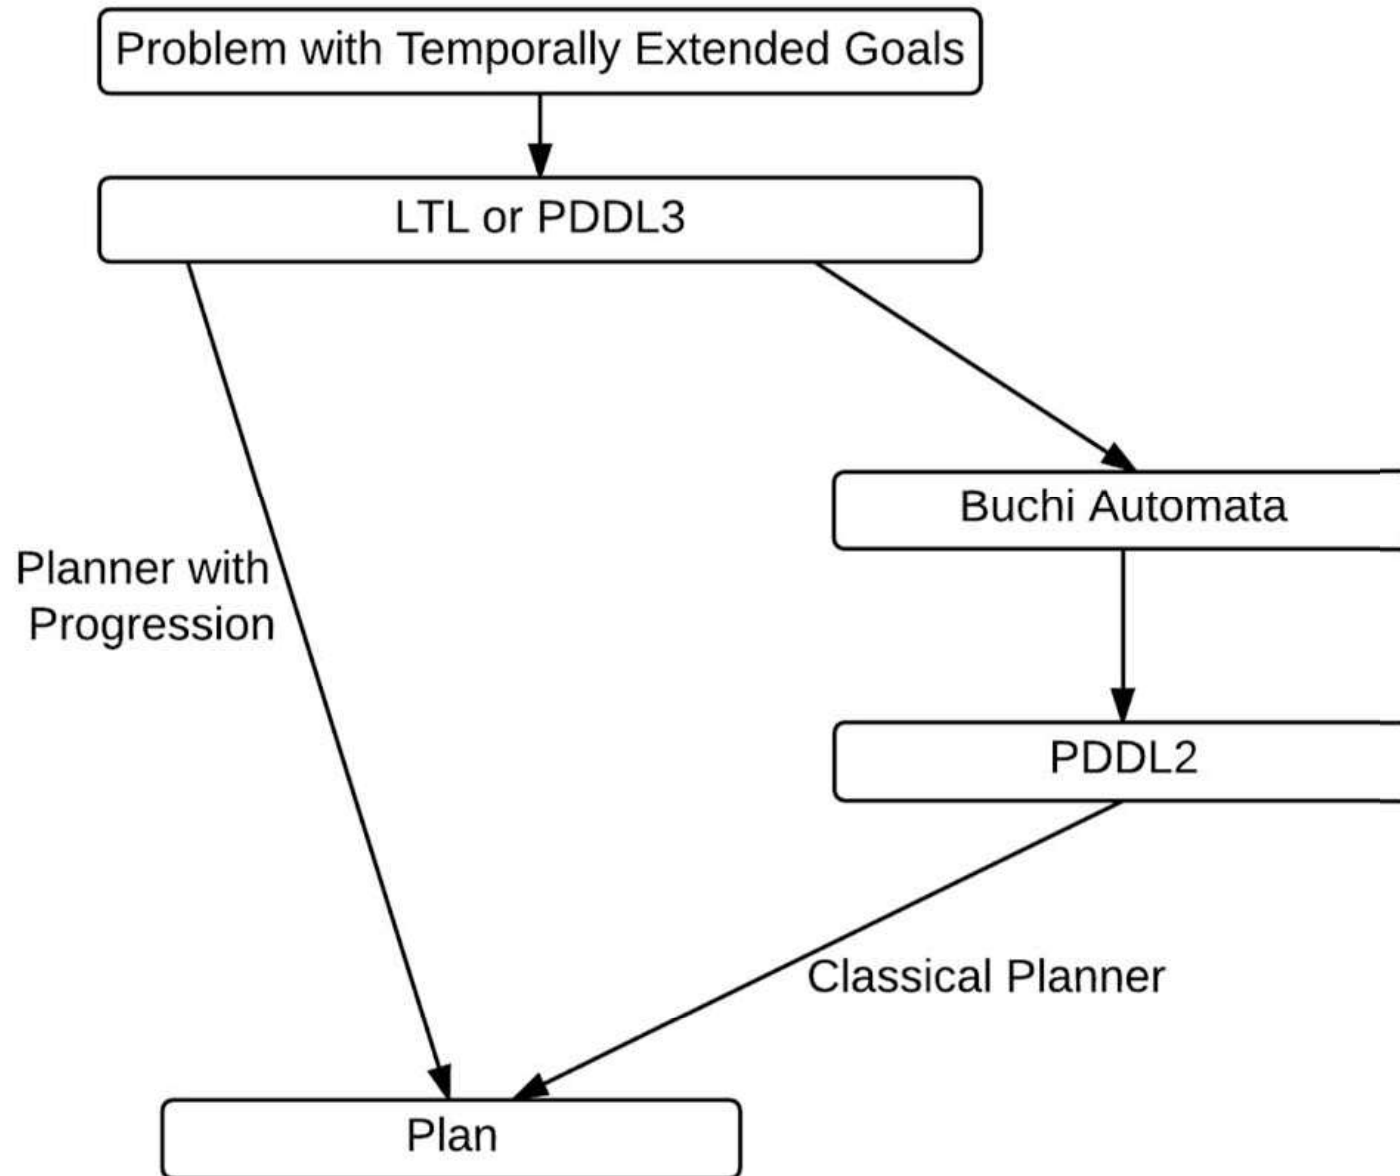
$$(g \textbf{U} r) \wedge (r \rightarrow \textbf{G}r)$$



```
(:goal
  (and
    (always-until (turn green) (turn
red) )
    (implies (turn red) (always (turn red)))
  )
)
```



Application to Planning





Büchi Automata

Büchi Automata - *extension of finite automaton to infinite inputs (words)*

A **Büchi automaton** is 5-tuple $\langle S, s_0, T, F, \Sigma \rangle$



- S is a finite set of **states**
- $s_0 \in S$ is an **initial state**
- $T \subseteq S \times \Sigma \rightarrow S$ is a **transition relation**
- $F \subseteq S$ is a set of **accepting states**
- Σ is a finite set of **symbols** ('alphabet')

An **infinite sequence of states** is accepted iff it visits the **accepting state(s) infinitely often**



Example Büchi Automata

Example: Model a clock



Accepted words:


TickTockTockTickTockTickTickTickTock...

TockTickTockTickTickTockTockTickTock...



Example Büchi Automata

Example: Model a clock




Accepted words:

TockTickTickTickTickTickTick...



Example Büchi Automata

Example: Model a clock




Accepted words:

TockTickTockTickTockTickTockTick...



LTL to Büchi Automata



neXt?


Future/Eventually?

Globally?



LTL to Büchi Automata


Future - $Fp \equiv \text{True} \cup p$



Accepted word: $\neg p \neg p \neg p p p \neg p \dots$

Sequence of states: $s_0 s_0 s_0 s_1 s_1 s_1 \dots$

Globally - $Gp \equiv \neg F\neg p$



Accepted word: $p p p p p \dots$

Sequence of states: $s_0 s_0 s_0 s_0 s_0 \dots$



LTL to Büchi Algorithm

N – Node object

N.curr – LTL formulas to be processed

N.old – LTL formulas already processed

N.next – LTL formulas to be processed in next node

N.incoming – Incoming transitions from predecessor nodes

N_s – List of processed Nodes

N_i – Arbitrary node from N_s

expand (N, N_s)

if N.curr is empty

if N.curr = N_i .curr

Append N.curr to N_i .curr

else

Append N to N_s

Create new node N_{new} with $N_{new}.curr = N.next$

expand (N_{new} , N_s)

else

Remove an LTL formula f from N.curr and append to N.old

Perform **Progression** on f

Call **expand** on result of **Progression**

The result of this algorithm is a generalized Buchi automata which is then transformed into a simple Buchi automata.



Progression Algorithm

progress(f, N, $\Delta t = 1$) *# Δt is time between successive states*

if f contains no temporal qualities:

if N.curr entails f:

$f' = \text{True}$

else

$f' = \text{False}$

if $f = f_1 \wedge f_2$:

progress(f_1 , N, Δt) \wedge **progress**(f_2 , N, Δt)

if $f = \mathbf{X}f_1$:

N.next.append(f_1)

if $f = f_1 \mathbf{U}_{[a,b]} f_2$:

$[a,b]$ is a time interval that could be infinite

if $b < a$:

$f' = \text{False}$

else if $0 \in [a,b]$:

progress(f_2 , N, Δt) \vee (**progress**(f_1 , N, Δt) \wedge N.next.append($f_1 \mathbf{U}_{[a,b] - \Delta t}$ f))

else


progress(f_1 , N, Δt) \wedge N.next.append($f_1 \mathbf{U}_{[a,b] - \Delta t}$ f)



Büchi Automata to PDDL2

Büchi states are not equivalent to PDDL2 states. Consider:

FutureGlobally - FGp



Two ways to transform temporally extended goals to PDDL2:

- Create new actions that encapsulate the allowable transitions in each state
- Introduce derived predicates
 - Do not depend on the actions
 - Used to determine which state the planner is in
 - Goal of the planner is to move from initial state to any accepting state



Planning with Preferences



Preference Based Planning

Classical Planning Problem

$$problem := (S, s_0, A, G)$$

S - set of states *s*₀ - initial state *A* - set of operators *G* - set of goal states

Preference-based Planning Problem

$$problem := (S, s_0, A, G, R)$$

R is a partial or total relation expressing preferences (\preceq)
between plans

**Preferences express properties of the plan
that are desired but not required**



Preference Expression Languages

- **Quantitative** - *assign numeric values to plans to compare them*
 - Markov Decision Processes (MDP's)
 - Find preferred policy using a reward function over conditional plans
 - PDDL3
 - Preferences expressed through reward function based on satisfying/violating logical formulas on the plan
- **Qualitative** - *relations compare plans based on properties of the plans that need not be numeric*
 - Ranked Knowledge Bases
 - Plan properties are ranked with preferred formulas ranked higher
 - Temporally Extended Preferences
 - Use LTL to express plan properties that are then ranked

Quantitative languages imply total comparability while qualitative languages may allow incomparability



Syntax for modeling preferences:

`(preference [name] <GD>)` - *label for fluents that represent preferences*

`is-violated` - *function that returns the number of times the preference was not satisfied in the plan*

Example:

Traffic light is green until it turns red

```
(preference gUr  
  (always-until (turn green) (turn  
red) ) )
```

Plan tries to not violate any preferred fluents

```
(metric minimize (is-violated gUr) )
```



LPP Language Overview

- LPP is a quantitative language to express temporal preferences for planning
 - Preferences between different temporal goals can be expressed along with the strength of preference
 - i.e. Goal A is preferred twice as much as Goal B
- LPP is an extension of an older language PP
- Preference formulas in LPP are constructed hierarchically

See Bienvenu, Meghyn, Christian Fritz, and Sheila A. McIlraith. "Planning with Qualitative Temporal Preferences." KR 6 (2006): 134-144.



Constructing a Preference Formula

Basic Desire Formula (BDFs)

express temporally extended propositions

- At some point, will cook
 - $b_1 = \mathbf{F}(\text{cook})$
- At some point, will order takeout
 - $b_2 = \mathbf{F}(\text{orderTakeout})$
- At some point, will eat spaghetti
 - $b_3 = \mathbf{F}(\text{eatSpaghetti})$
- At some point, will eat pizza
 - $b_4 = \mathbf{F}(\text{eatPizza})$



Constructing a Preference Formula

Atomic Preference Formulas (APFs) *express preferences between BDFs*

- In this example, weights associated with each BDF define preferences
 - Lower weight is preferred
- Prefer to cook over ordering takeout
 - $a_1 = b_1[0.2] \gg b_2[0.4]$
- Prefer eating spaghetti over eating pizza
 - $a_2 = b_3[0.3] \gg b_4[0.9]$



Constructing a Preference Formula

General Preference Formulas (GPFs)

allow conjunctions or disjunctions of APFs or qualification of BDFs with conditionals

- Satisfy the most preferred option among the APFs (satisfy APF with lowest weight)
 $-g_1 = a_1 \mid a_2$
- Choose the most preferred option that satisfies both APFs (minimize the maximum weight across both APFs)
 $-g_2 = a_1 \ \& \ a_2$



Constructing a Preference Formula

Aggregated Preferences Formulas (APFs)

define the order in which preferences should be relaxed

- Prefer that if both g_1 and g_2 from previous slide can't be met, that g_2 from previous slide is met
$$-g_1 \wedge g_2 \preceq g_2 \preceq g_1$$
- Situations that aren't distinguished any other way can be sorted lexicographically (alphabetically)



LPP Formula Hierarchy Review

- **Basic Desire Formula (BDF)**
 - Express temporally extended propositions
- **Atomic Preference Formula (APF)**
 - Express preferences between BDFs
- **General Preference Formula (GPF)**
 - Allow conjunctions or disjunctions of APFs or qualification of BDFs with conditionals
- **Aggregated Preference Formula (APF)**
 - Define the order in which preferences should be relaxed



References

- Gerevini, A., and D. Long. *Plan constraints and preferences in PDDL3: The language of the fifth international planning competition*. University of Brescia. Italy, Tech. Rep, 2005.
- Patrizeni, Fabio, et al. "Computing infinite plans for LTL goals using a classical planner." *Proceedings of the Twenty-Second International Joint Conference on Artificial Intelligence; July 16-22, 2011; Barcelona. Menlo Park, California: AAAI Press; 2011. p. 2003-2008.*. Association for the Advancement of Artificial Intelligence (AAAI), 2011.
- Baier, Jorge A., and Sheila A. McIlraith. "Planning with Temporally Extended Goals Using Heuristic Search." *ICAPS*. 2006.
- Bacchus, Fahiem, and Froduald Kabanza. "Planning for temporally extended goals." *Annals of Mathematics and Artificial Intelligence* 22.1-2 (1998): 5-27.
- Baier, Jorge A., and Sheila A. McIlraith. "Planning with first-order temporally extended goals using heuristic search." *Proceedings of the National Conference on Artificial Intelligence*. Vol. 21. No. 1. Menlo Park, CA; Cambridge, MA; London; AAAI Press; MIT Press; 1999, 2006.
- Baier, Jorge A., and Sheila A. McIlraith. "Planning with preferences." *AI Magazine* 29.4 (2009): 25.
- Bienvenu, Meghyn, Christian Fritz, and Sheila A. McIlraith. "Planning with Qualitative Temporal Preferences." *KR* 6 (2006): 134-144.
- Gerth, Rob, et al. "Simple on-the-fly automatic verification of linear temporal logic." *Protocol Specification, Testing and Verification XV*. Springer US, 1996. 3-18.



Appendix



- PPLAN
 - implemented by Meghyn Bienvenu, Christian Fritz, and Sheila A. McIlraith
- Solves planning problems with preferences expressed in LPP via bounded best-first search forward chaining planner
 - use of *progression* efficiently evaluates how well partial plans satisfy Φ (a general preference formula)
 - use of admissible *evaluation function* ensures best-first search is optimal